



# Deep Learning for Visual Computing

## Machine Learning for Image Classification

Christopher Pramerdorfer  
Computer Vision Lab, TU Wien

# Topics

Machine Learning recap

Image classification using Nearest Neighbor classifiers

Validation sets and hyperparameter optimization

# Motivation

We've seen that

- ▶ We humans are great image classifiers
- ▶ But we cannot formally describe how we do it
- ▶ So we can't write algorithms for reliable classification

This applies to most vision problems

- ▶ Reason we need Machine and Deep Learning

# Motivation

Machine Learning (ML) algorithms are able to learn from data

- ▶ Algorithm performance improves with experience

In this framework we

- ▶ Select a suitable ML algorithm
- ▶ Collect data for algorithm training and evaluation
- ▶ Monitor the training progress

# ML for Image Classification

Recall that we want to build an image classifier

- ▶ Should support the classes {dog, cat}
- ▶ Using the CIFAR-10 dataset (6000 images per class)



Image from [cs.toronto.edu](http://cs.toronto.edu)

# ML for Image Classification

In this context

- ▶ We show (image, class) pairs to the chosen ML algorithm
- ▶ Algorithm learns to predict class of **unseen** samples

The information we are interested in is called **label**

- ▶ In our case the label value is *cat* or *dog*
- ▶ Finite number of label values, hence **classification** problem
- ▶ Labels used during training, hence **supervised learning**




# ML for Image Classification

## First Try

Represent images as points in  $D$ -dimensional **input space**

- ▶ Stack image rows/columns to obtain vector  $\mathbf{x}$
- ▶  $D$  is usually large (CIFAR-10 :  $D = 32 \cdot 32 \cdot 3 = 3072$ )





# ML for Image Classification

## First Try

Assuming  $D = 2$  and three classes (colors)

- Can you think of a simple approach for classification?




Image adapted from [cs231n.github.io](https://github.com/cs231n)

# ML for Image Classification

## First Try

### Nearest Neighbor classifier

- ▶ Compute distance from  $\times$  to all training samples
- ▶ Assign label of closest sample




Image adapted from [cs231n.github.io](https://cs231n.github.io)

# ML for Image Classification

## First Try

Need **distance measure** between two images  $\mathbf{x}_1, \mathbf{x}_2$

$$\text{L1 distance} : \sum_{d=1}^D |x_1^d - x_2^d|$$

$$\text{L2 (Euclidean) distance} : \sqrt{\sum_{d=1}^D (x_1^d - x_2^d)^2}$$

# ML for Image Classification

## First Try

Resulting **decision regions** and **decision boundaries**

- ▶ Voronoi tessellation of input space





Image from [cs231n.github.io](https://cs231n.github.io)

# ML for Image Classification

## First Try

Performance (on full CIFAR-10 dataset)

- ▶ Obvious errors and challenges the classifier is not robust to
- ▶ But performance likely better than a “manual” method






# ML for Image Classification

## Second Try

How can we improve the performance?



# ML for Image Classification

## Second Try

### $k$ Nearest Neighbors classifier

- ▶ Find labels of  $k$  closest training samples
- ▶ Assign most frequent label




Image adapted from [cs231n.github.io](https://cs231n.github.io)



# ML for Image Classification

## Second Try

Results in smoother decision boundaries

- Visualization suggest better performance on unseen samples




Image adapted from [cs231n.github.io](https://cs231n.github.io)

# ML for Image Classification

## Second Try

We want to know how certain the classifier is

- ▶ Allows us to react differently on this basis

We can adapt the classifier accordingly

- ▶ Predict **class scores**  $\mathbf{w} \in \mathbb{R}^C$  ( $C$  is number of classes)
- ▶  $w_c$  is frequency of label  $c$  among  $k$  closest samples
- ▶ In example  $\mathbf{w} = (1, 0, 4)$  assuming red, green, blue order

# ML for Image Classification

## Second Try

Class scores computed this way are not optimal

- ▶ Absolute values that depend on  $k$
- ▶ Other classifiers will compute scores differently

We want  $\mathbf{w}$  to be a valid probability mass function

- ▶  $w_c \geq 0$  for all  $c$  and  $\sum_c w_c = 1$

Most popular function for this purpose is **softmax**

$$\text{softmax}_c(\mathbf{w}) = \frac{\exp(w_c)}{\sum_c \exp(w_c)}$$

We obtain  $\text{softmax}((1, 0, 4)) \approx (0.05, 0.02, 0.93)$

- ▶ Largest value is emphasized, small ones suppressed
- ▶ Softmax is not scale invariant

# ML for Image Classification

## Second Try

How should we select  $k$  for cat vs. dog classification?

- ▶ Has a large impact on performance
- ▶ Different  $k$  work best depending on the data
- ▶ Cannot visualize the 3072-dimensional input space

$k$  is a **hyperparameter**

- ▶ Set manually, controls algorithm behavior
- ▶ Chosen experimentally or based on domain knowledge

# ML for Image Classification

## Second Try

To find a good  $k$  for our problem, we

- ▶ Sample  $k$  from a reasonable range (what's reasonable?)
- ▶ Quantify the algorithm performance on a **validation set**
- ▶ Chose the  $k$  with the highest performance


Why not use the training or test sets?

# ML for Image Classification

## Second Try

We thus need three **disjoint** datasets

- ▶ Training set for classifier training (duh)
- ▶ Validation set for hyperparameter selection
- ▶ Test set for final estimate of performance in practice




# ML for Image Classification

## Second Try

Test set used to estimate the performance in practice

- ▶ No longer valid if used during training/validation
- ▶ Use only once at the very end





# ML for Image Classification

## Second Try

Need suitable **performance measure** to quantify performance

**Accuracy** is popular for classification

- ▶ Let algorithm predict labels for dataset
- ▶ Compare predictions and true (**ground truth**) labels
- ▶ Accuracy is fraction of correctly classified samples

Measure 1 – accuracy is called **error rate**

# ML for Image Classification

## Second Try

We generally cannot test all hyperparameter combinations

- ▶ Testing one combination can take long
- ▶ Number blows up if we have several hyperparameters
- ▶ Large and/or continuous intervals

Use an approximative search strategy

- ▶ Grid search
- ▶ Random search

# ML for Image Classification

## Second Try

Given  $H$  hyperparameters with search intervals  $I_1 \cdots I_H$

### Grid search

- ▶ Sample uniformly from all  $I_h$
- ▶ Test all combinations

### Random search

- ▶ For  $j$  iterations, sample randomly from all  $I_h$
- ▶ Test sample combination

# ML for Image Classification

## Second Try

Random search usually works better

- Wastes less time on unimportant hyperparameters




Image from [1]

# ML for Image Classification

## Second Try

$k$  Nearest Neighbor classifier performs better but still not great

- ▶ More improvements in next lecture

- [1] James Bergstra and Yoshua Bengio. *Random search for hyper-parameter optimization*. Journal of Machine Learning Research (2012).