



# Deep Learning for Visual Computing

## Loss Functions, Iterative Optimization

Christopher Pramerdorfer  
Computer Vision Lab, TU Wien

# Topics

How good is our classifier?

- ▶ Cross-entropy loss

How should we adapt the parameters?

- ▶ Iterative optimization
- ▶ Gradient Descent

# Training Parametric Models

In parametric models  $f$  depends on parameters  $\theta$

- ▶ We write  $\mathbf{w} = f(\mathbf{x}; \theta)$
- ▶ Training entails finding good parameters

For training any parametric model we need

- ▶ A loss function
- ▶ An optimization algorithm

# Loss Functions

A **loss function**  $L(\theta)$  (or **cost** or **objective** function)

- ▶ Measures performance of  $f(\cdot; \theta)$  (lower loss is better)
- ▶ On some (training) dataset  $\mathcal{D} = \{(\mathbf{x}_s, \mathbf{w}_s)\}_{s=1}^S$
- ▶ With respect to parameters  $\theta$

Choice of  $L$  depends on task

- ▶ Most popular classification loss is cross-entropy

# Loss Functions

## Cross-Entropy

Given two probability mass functions  $\mathbf{u}$  and  $\mathbf{v}$  in  $\mathbb{R}^T$

►  $\mathbf{u} = (u_1, \dots, u_T)$  and  $\mathbf{v} = (v_1, \dots, v_T)$

The **cross-entropy** between  $\mathbf{u}$  and  $\mathbf{v}$  is


$$H(\mathbf{u}, \mathbf{v}) = - \sum_{t=1}^T u_t \ln v_t$$

# Loss Functions

## Cross-Entropy

Example with  $T = 2$  and  $u_1 = 1$

- ▶ The more different  $\mathbf{u}$  and  $\mathbf{v}$  the higher  $H$
- ▶  $H$  measures the dissimilarity between  $\mathbf{u}$  and  $\mathbf{v}$




# Loss Functions

## Cross-Entropy

Note that  $H$  can reach 0 only if any  $u_t = 1$

- In general  $H(\mathbf{u}, \mathbf{v}) = \text{entropy of } \mathbf{u}$  if  $\mathbf{u} = \mathbf{v}$



# Loss Functions

## Cross-Entropy Loss

To utilize the cross-entropy for classifier training we

- ▶ Let  $\mathbf{u}$  encode the ground-truth label,  $u_c = 1$
- ▶ Let  $\mathbf{v}$  be the predicted softmax class scores

$H$  measures how dissimilar true and predicted probabilities are

- ▶ How well the classifier performs on a single sample



# Loss Functions

## Cross-Entropy Loss

On this basis we calculate the **cross-entropy loss** on  $\mathcal{D}$  as

$$L(\theta) = \frac{1}{S} \sum_{s=1}^S H(\mathbf{w}_s, \text{softmax}(f(\mathbf{x}_s; \theta)))$$

Average cross-entropy over some dataset  $\mathcal{D}$

- We will use (subsets of) the training set as  $\mathcal{D}$

# Loss Functions

## Cross-Entropy Loss

Models trained with this loss are called **softmax classifiers**

- ▶ Also called **logistic regression** if  $T = 2$

Classifiers learn to predict probabilities per class label

- ▶ In theory predictions are reliable probability estimates
- ▶ In practice DL classifiers are often overconfident



# Gradient Descent

## Motivation

We now know how to compute  $L(\theta)$  for classification

Need a way to minimize  $L(\theta)$

- ▶ Maximizes the training set classification performance
- ▶ And hopefully also validation/test performance (more later)

$L(\theta)$  is not linear in  $\theta$

- ▶ Need a **nonlinear optimization** algorithm
- ▶ **Gradient Descent** is popular choice in Deep Learning (DL)

# Gradient Descent

## Introduction


Assume terrain corresponds to  $L(\theta)$  with  $\dim(\theta) = 2$



# Gradient Descent

## Introduction

How do I get from location  $\theta$  to location of minimum  $\hat{\theta}$ ?



# Gradient Descent

## Introduction

Without actually seeing  $L(\theta)$ ?



# Gradient Descent

## Introduction

Feel slope with feet, step in direction that feels steepest

- ▶ Again and again until ground feels flat





# Gradient Descent

## Definition

### Iterative Optimization algorithm

In every iteration we

- ▶ Compute gradient  $\theta' = \nabla L(\theta)$
- ▶ Update parameters  $\theta = \theta - \alpha \theta'$

Hyperparameter  $\alpha > 0$  is called **learning rate**

- ▶ Final **step size** is  $\alpha \|\theta'\|$

# Gradient Descent

## Gradients

Let  $f(x_1, \dots, x_n)$  be a differentiable, real-valued function

The **partial derivative**  $f_{x_i}$  of  $f$  with respect to  $x_i$

- ▶ Is also a real-valued function  $f_{x_i}(x_1, \dots, x_n)$

$f_{x_i}(\mathbf{x})$  encodes

- ▶ How fast  $f$  changes with argument  $x_i$
- ▶ At some location  $\mathbf{x}$

# Gradient Descent

## Gradients

**Gradient**  $\nabla f$  is vector of all partial derivatives of  $f$

- ▶  $\nabla f = (f_{x_1}, \dots, f_{x_n})$
- ▶ Vector-valued function  $\mathbb{R}^n \mapsto \mathbb{R}^n$

$\nabla f(\mathbf{x}) = (f_{x_1}(\mathbf{x}), \dots, f_{x_n}(\mathbf{x}))$  encodes

- ▶ How fast  $f$  changes with all arguments  $x_1 \cdots x_n$
- ▶ At some location  $\mathbf{x}$

# Gradient Descent

## Gradients




Image adapted from [wikimedia.org](https://commons.wikimedia.org/wiki/File:Gradient_descent_3D.png)

# Gradient Descent

## Gradients

$\nabla f(\mathbf{x})$  specifies how  $f$  changes locally at  $\mathbf{x}$

- ▶ Points in direction of greatest increase
- ▶ Norm equals magnitude of increase

Exactly what we need to minimize  $L$

- ▶ Compute direction of greatest increase  $\nabla L(\boldsymbol{\theta})$
- ▶ Move in the opposite direction

# Gradient Descent

## Gradients

We stop if  $\nabla L(\theta) \approx \mathbf{0}$  (if norm is close to 0)

- ▶ No information where to go next
- ▶  $L$  is flat at current location
- ▶ The case if we are at  $\hat{\theta}$  (but not only then)

# Gradient Descent

## Remarks

Simple and general algorithm

- ▶ Requires only that  $f$  is differentiable, real-valued
- ▶ Efficient (requires only first derivatives)

Several (possible) limitations

- ▶ Performs poorly for many  $f$
- ▶ But works remarkable well with DL models

# Gradient Descent

## Limitations – Critical Points and Local Minima

Algorithm stops if  $\nabla L(\theta) \approx 0$

- ▶ Applies to all **critical points**, not only minimum
- ▶ Should stop only at minimum




Image from [1]



# Gradient Descent

## Limitations – Critical Points and Local Minima

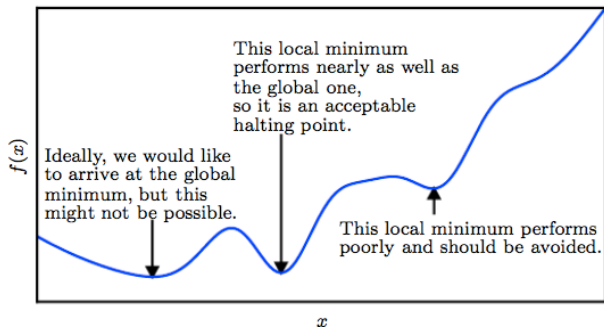


Image from [1]

# Gradient Descent

## Limitations – Critical Points and Local Minima

Algorithm stops at first minimum as  $\nabla L(\theta) \approx 0$

- ▶ But  $L$  generally has several **local minima**
- ▶ Algorithm usually finds only a local minimum

For loss functions of DL models evaluated on minibatches (below)

- ▶ Local minima are usually close to **global minimum**
- ▶ Optimization does not come close to critical points

# Gradient Descent

## Limitations – Poorly Conditioned Hessian

Very different curvature in different directions (canyon-like)




Image from [1]

# Gradient Descent

Limitations – Poorly Conditioned Hessian

Gradient descent wastes time jumping between canyon walls




Image from [1]

# Gradient Descent

## Momentum

Momentum improves speed of convergence by

- ▶ Dampening oscillations (previous slide)
- ▶ Increasing step size dynamically

Use exponential moving average of gradients for direction  $\mathbf{v}$

- ▶ Influence of older gradients decays exponentially

# Gradient Descent

## Momentum

Iteration of gradient descent with momentum

- ▶ Update velocity  $\mathbf{v} = \beta \mathbf{v} - \alpha \nabla L(\boldsymbol{\theta})$
- ▶ Update parameters  $\boldsymbol{\theta} = \boldsymbol{\theta} + \mathbf{v}$

Hyperparameter  $\beta \in [0, 1)$  called **momentum**

- ▶ Defines decay speed and maximum step size

# Gradient Descent

## Momentum

$\mathbf{v}$  builds up momentum if successive gradients are similar

- ▶ Improves speed of convergence

Maximum step size is  $\alpha \|\mathbf{g}\| / (1 - \beta)$

- ▶ Assuming the gradient is always  $\mathbf{g}$
- ▶ At  $\beta = 0.9$  maximum increase by factor of 10

# Gradient Descent

## Momentum

Red is path, black are steepest descent directions




Image from [1]



# Gradient Descent

## Nesterov Momentum

Evaluate gradient at  $\theta + \mathbf{v}$  instead of  $\theta$

Iteration of gradient descent with **Nesterov momentum**

- ▶ Update velocity  $\mathbf{v} = \beta \mathbf{v} - \alpha \nabla L(\theta + \mathbf{v})$
- ▶ Update parameters  $\theta = \theta + \mathbf{v}$

Often works better than standard momentum

# Gradient Descent

## Stochastic Gradient Descent

Goal is to minimize  $L(\theta)$  as measured on training data

Obvious choice is to use whole training set

- ▶ Called **Batch Gradient Descent**
- ▶ Time complexity per iteration increases linearly with  $S$
- ▶ Problematic if  $S$  is large (need many iterations)

# Gradient Descent

## Stochastic Gradient Descent

To solve this problem we

- ▶ Process the whole training set
- ▶ In **minibatches** of size  $S$  (one per iteration)

Possible because gradient is an expectation

- ▶ Can estimate training set loss on subset
- ▶ Also applies for the gradient

One full run through the training set is called an **epoch**

- ▶ Usually training takes many epochs

# Gradient Descent

## Stochastic Gradient Descent

Resulting algorithm called **Minibatch Gradient Descent**

- ▶ Or **Stochastic Gradient Descent (SGD)** if  $S = 1$
- ▶ In practice often called SGD even if  $S > 1$

Time for single iteration is now independent of dataset size

In DL  $S$  varies between 1 and a few hundred samples

- ▶ Most common are 64, 128, 256
- ▶  $2^n$  for efficiency (data parallelism)

# Gradient Descent

## Stochastic Gradient Descent

Decreasing  $S$  also decreases

- ▶ Computation time per iteration
- ▶ Memory required on GPU (minibatch processed as whole)
- ▶ Accuracy of the gradient estimate

Decreasing  $S$  causes more noisy gradient estimates

- ▶ Gives Gradient Descent ability to escape local minima [2]

# Gradient Descent

## Stochastic Gradient Descent

Important to sample minibatches randomly

- ▶ To break (possible) ordering in dataset

Standard approach in practice

- ▶ Shuffle training set once or before every epoch
- ▶ Process sequentially in minibatches

# Gradient Descent

## Alternatives

Many alternatives

- ▶ Adagrad, RMSProp, Adam, ...
- ▶ Advantage of not having to choose the learning rate

Overall SGD with Nesterov momentum is the best choice

- ▶ Setting  $\beta = 0.9$  is usually fine

# Gradient Descent

## Alternatives

Path finding comparison on challenging  $f$

- Different learning rates, so speed not comparable




Image from [github.com](https://github.com)



# Bibliography

- [1] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. 2016.
- [2] Robert Kleinberg, Yuanzhi Li, and Yang Yuan. *An Alternative View: When Does SGD Escape Local Minima?* CoRR abs/1802.06175 (2018). URL: <http://arxiv.org/abs/1802.06175>.